

HW5

Запуск

1) Композ файл с монтированием директории:

```
services:  
  spark:  
    build:  
      context: .  
      dockerfile: Dockerfile  
    image: spark-python  
    container_name: spark-python  
    volumes:  
      - ./app:/app
```

2) Докерфайл из образа apache/spark:3.5.0 с установкой python и pyspark:

FROM apache/spark:3.5.0

```
USER root

RUN apt-get update && \
    apt-get install -y python3 python3-pip && \
    ln -sf /usr/bin/python3 /usr/bin/python && \
    pip3 install --no-cache-dir pyspark && \
    apt-get clean && rm -rf /var/lib/apt/lists/*
```

WORKDIR /app

```
CMD ["bash", "run_test.sh"]
```

3) запуск spark приложения docker compose up

Spark приложение

1) `test_spark_df.py`: Обработка csv, агрегация и усреднение по колонке, вывод с сортировкой.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import avg, count

spark = SparkSession.builder.appName("CSVStats").getOrCreate()

df = spark.read.csv("./data/flights.csv", header=True, inferSchema=True)

df.groupBy("DayofMonth").agg(
    count("*").alias("rows"),
```

```

    avg("DepDelay").alias("avg_DepDelay")
).orderBy("DayofMonth").show(
    n=10,
    truncate=False,
    vertical=False
)

spark.stop()

```

Вывод:

DayofMonth	rows	avg_DepDelay
1	84636	8.868507490902216
2	89760	10.928219696969697
3	90172	9.696346981324579
4	84758	6.338304348852025
5	86426	6.2484900377201305
6	87702	7.087899933866958
7	88011	10.064753269477679
8	89019	10.662768622428919
9	91412	11.449831531965168
10	90025	16.279622327131353

- 2) test_spark_df.py: Обработка таблиц через SQL синтаксис, объединение таблиц, агрегация и усреднение по колонке, вывод с сортировкой

```

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("SparkSQLJoin").getOrCreate()

flights = spark.read.csv("./data/flights.csv", header=True, inferSchema=True)
carriers = spark.read.csv("./data/airports.csv", header=True, inferSchema=True)

flights.createOrReplaceTempView("flights")
carriers.createOrReplaceTempView("airports")

spark.sql("""
    SELECT
        a1.name AS origin_airport,
        a2.name AS destination_airport,
        ROUND(AVG(f.DepDelay),2) AS avg_dep_delay
    FROM flights f
    JOIN airports a1 ON f.OriginAirportID = a1.airport_id
    JOIN airports a2 ON f.DestAirportID = a2.airport_id
    WHERE f.DepDelay IS NOT NULL
""")

```

```

        GROUP BY a1.name, a2.name
        HAVING AVG(f.DepDelay) > 30
        ORDER BY avg_dep_delay DESC;
    """).show(
        n=10,
        truncate=False,
        vertical=False
    )

spark.stop()

```

Выход:

spark-python	origin_airport	destination_airport
spark-python	Pittsburgh International	Richmond International
spark-python	Pittsburgh International	Raleigh-Durham International
spark-python	Los Angeles International	Eppley Airfield
spark-python	Cincinnati/Northern Kentucky International	Cleveland-Hopkins International
spark-python	Seattle/Tacoma International	Miami International
spark-python	Chicago Midway International	Ontario International
spark-python	Fort Lauderdale-Hollywood International	Richmond International
spark-python	Chicago Midway International	San Francisco International
spark-python	Norfolk International	Minneapolis-St Paul International
spark-python	Metropolitan Oakland International	Logan International
spark-python	only showing top 10 rows	